

OpenCV 4.x Cheat Sheet (Python version)

A summary of: <https://docs.opencv.org/master/>

I/O

<code>i = imread("name.png")</code>	Loads image as BGR (if grayscale, B=G=R)
<code>i = imread("name.png", IMREAD_UNCHANGED)</code>	Loads image as is (inc. transparency if available)
<code>i = imread("name.png", IMREAD_GRAYSCALE)</code>	Loads image as grayscale
<code>imshow("Title", i)</code>	Displays image I
<code>imwrite("name.png", i)</code>	Saves image I
<code>waitKey(500)</code>	Wait 0.5 seconds for keypress (0 waits forever)
<code>destroyAllWindows()</code>	Releases and closes all windows

Color/Intensity

<code>i_gray = cvtColor(i, COLOR_BGR2GRAY)</code>	BGR to gray conversion
<code>i_rgb = cvtColor(i, COLOR_BGR2RGB)</code>	BGR to RGB (useful for <code>matplotlib</code>)
<code>i = cvtColor(i, COLOR_GRAY2RGB)</code>	Converts grayscale to RGB (R=G=B)
<code>i = equalizeHist(i)</code>	Histogram equalization
<code>i = normalize(i, None, 0, 255, NORM_MINMAX, CV_8U)</code>	Normalizes I between 0 and 255
<code>i = normalize(i, None, 0, 1, NORM_MINMAX, CV_32F)</code>	Normalizes I between 0 and 1

Other useful color spaces

<code>COLOR_BGR2HSV</code>	BGR to HSV (Hue, Saturation, Value)
<code>COLOR_BGR2LAB</code>	BGR to Lab (Lightness, Green/Magenta, Blue/Yellow)
<code>COLOR_BGR2LUV</code>	BGR to Luv (\approx Lab, but different normalization)
<code>COLOR_BGR2YCrCb</code>	BGR to YCrCb (Luma, Blue-Luma, Red-Luma)

Channel manipulation

<code>b, g, r = split(i)</code>	Splits the image I into channels
<code>b, g, r, a = split(i)</code>	Same as above, but I has alpha channel
<code>i = merge((b, g, r))</code>	Merges channels into image

Arithmetic operations

<code>i = add(i1, i2)</code>	$\min(I_1 + I_2, 255)$, i.e. saturated addition if <code>uint8</code>
<code>i = addWeighted(i1, alpha, i2, beta, gamma)</code>	$\min(\alpha I_1 + \beta I_2 + \gamma, 255)$, i.e. image blending
<code>i = subtract(i1, i2)</code>	$\max(I_1 - I_2, 0)$, i.e. saturated subtraction if <code>uint8</code>
<code>i = absdiff(i1, i2)</code>	$ I_1 - I_2 $, i.e. absolute difference

Note: one of the images can be replaced by a scalar.

Logical operations

<code>i = bitwise_not(i)</code>	Inverts every bit in I (e.g. mask inversion)
<code>i = bitwise_and(i1, i2)</code>	Logical <i>and</i> between I_1 and I_2 (e.g. mask image)
<code>i = bitwise_or(i1, i2)</code>	Logical <i>or</i> between I_1 and I_2 (e.g. merge 2 masks)
<code>i = bitwise_xor(i1, i2)</code>	Exclusive <i>or</i> between I_1 and I_2

Statistics

<code>mB, mG, mR, mA = mean(i)</code>	Average of each channel (i.e. BGRA)
<code>ms, sds = meanStdDev(i)</code>	Mean and SDev p/channel (3 or 4 rows each)
<code>h = calcHist([i], [c], None, [256], [0,256])</code>	Histogram of channel c , no mask, 256 bins (0-255)
<code>h = calcHist([i], [0,1], None, [256,256], [0,256, 0,256])</code>	2D histogram using channels 0 and 1, with “resolution” 256 in each dimension

Filtering

<code>i = blur(i, (5, 5))</code>	Filters I with 5×5 box filter (i.e. average filter)
<code>i = GaussianBlur(i, (5,5), sigmaX=0, sigmaY=0)</code>	Filters I with 5×5 Gaussian; auto σ s; (I is <code>float</code>)
<code>i = GaussianBlur(i, None, sigmaX=2, sigmaY=2)</code>	Blurs, auto kernel dimension
<code>i = filter2D(i, -1, k)</code>	Filters with 2D kernel using cross-correlation
<code>kx = getGaussianKernel(5, -1)</code>	1D Gaussian kernel with length 5 (auto StDev)
<code>i = sepFilter2D(i, -1, kx, ky)</code>	Filter using separable kernel (same output type)
<code>i = medianBlur(i, 3)</code>	Median filter with size=3 (size ≥ 3)
<code>i = bilateralFilter(i, -1, 10, 50)</code>	Bilateral filter with $\sigma_r = 10$, $\sigma_s = 50$, auto size

Borders

All filtering operations have parameter `borderType` which can be set to:

<code>BORDER_CONSTANT</code>	Pads with constant border (requires additional parameter <code>value</code>)
<code>BORDER_REPLICATE</code>	Replicates the first/last row and column onto the padding
<code>BORDER_REFLECT</code>	Reflects the image borders onto the padding
<code>BORDER_REFLECT_101</code>	Same as previous, but doesn't include the pixel at the border (the default)
<code>BORDER_WRAP</code>	Wraps around the image borders to build the padding

Borders can also be added with custom widths:

<code>i = copyMakeBorder(i, 2, 2, 3, 1, borderType=BORDER_WRAP)</code>	Widths: top, bottom, left, right
--	----------------------------------

Differential operators

<code>i_x = Sobel(i, CV_32F, 1, 0)</code>	Sobel in the x-direction: $I_x = \frac{\partial}{\partial x} I$
<code>i_y = Sobel(i, CV_32F, 0, 1)</code>	Sobel in the y-direction: $I_y = \frac{\partial}{\partial y} I$
<code>i_x, i_y = spatialGradient(i, 3)</code>	The gradient: ∇I (using 3×3 Sobel): needs <code>uint8</code> image
<code>m = magnitude(i_x, i_y)</code>	$\ \nabla I\ $; I_x, I_y must be float (for conversion, see <code>np.astype()</code>)
<code>m, d = cartToPolar(i_x, i_y)</code>	$\ \nabla I\ $; $\theta \in [0, 2\pi]$; <code>angleInDegrees=False</code> ; needs <code>float32</code> I_x, I_y
<code>l = Laplacian(i, CV_32F, ksize=5)</code>	ΔI , Laplacian with kernel size of 5

Geometric transforms

<code>i = resize(i, (width, height))</code>	Resizes image to <code>width</code> \times <code>height</code>
<code>i = resize(i, None, fx=0.2, fy=0.1)</code>	Scales image to 20% width and 10% height
<code>M = getRotationMatrix2D((xc, yc), deg, scale)</code>	Returns 2×3 rotation matrix M , arbitrary (x_c, y_c)
<code>M = getAffineTransform(pts1,pts2)</code>	Affine transform matrix M from 3 correspondences
<code>i = warpAffine(i, M, (cols,rows))</code>	Applies Affine transform M to I , output size=(<code>cols</code> , <code>rows</code>)
<code>M = getPerspectiveTransform(pts1,pts2)</code>	Perspective transform matrix M from 4 correspondences
<code>M, s = findHomography(pts1, pts2)</code>	Persp transf $m \times M$ from all $\gg 4$ corresps (Least squares)
<code>M, s = findHomography(pts1, pts2, RANSAC)</code>	Persp transf $m \times M$ from best $\gg 4$ corresps (RANSAC)
<code>i = warpPerspective(i, M, (cols, rows))</code>	Applies perspective transform M to image I

Interpolation methods

`resize`, `warpAffine` and `warpPerspective` use bilinear interpolation by default. It can be changed by parameter `interpolation` for `resize`, and `flags` for the others:

<code>flags=INTER_NEAREST</code>	Simplest, fastest (or <code>interpolation=INTER_NEAREST</code>)
<code>flags=INTER_LINEAR</code>	Bilinear interpolation: Default
<code>flags=INTER_CUBIC</code>	Bicubic interpolation

Segmentation

<code>_, i_t = threshold(i, t, 255, THRESH_BINARY)</code>	Manually thresholds image I given threshold level t
<code>t, i_t = threshold(i, 0, 255, THRESH_OTSU)</code>	Returns thresh level and thresholded image using Otsu
<code>i_t = adaptiveThreshold(i, 255, ADAPTIVE_THRESH_MEAN_C, THRESH_BINARY, b, c)</code>	Adaptive mean-c with block size b and constant c
<code>bp = calcBackProject([i_hsv], [0,1], h, [0,180, 0,256], 1)</code>	Back-projects histogram h onto the image <code>i_hsv</code> using only hue and saturation; no scaling (i.e. 1)
<code>cp, la, ct = kmeans(feats, k, None, crit, 10, KMEANS_RANDOM_CENTERS)</code>	Returns the labels <code>la</code> and centers <code>ct</code> of K clusters, best compactness <code>cp</code> out of 10; 1 feat/column

Features

e = Canny(i, tl, th)	Returns the Canny edges (e is binary)
l = HoughLines(e, 1, pi/180, 150)	Returns all $(\rho, \theta) \geq 150$ votes, Bin res: $\rho = 1$ pix, $\theta = 1$ deg
l = HoughLinesP(e, 1, pi/180, 150, None, 100, 20)	Probabilistic Hough, min length=100, max gap=20
c = HoughCircles(i, HOUGH_GRADIENT, 1, minDist=50, param1=200, param2=18, minRadius=20, maxRadius=60)	Returns all (x_c, y_c, r) with at least 18 votes, bin resolution=1, param1 is the t_h of Canny, and the centers must be at least 50 pixels away from each other
r = cornerHarris(i, 3, 5, 0.04)	Harris corners' R_s per pixel, window=3, Sobel=5, $\alpha = 0.04$
f = FastFeatureDetector_create()	Instantiates the Star feature detector
k = f.detect(i, None)	Detects keypoints on grayscale image I
i_k = drawKeypoints(i, k, None)	Draws keypoints k on color image I
d = xfeatures2d.BriefDescriptorExtractor_create()	Instantiates a BRIEF descriptor
k, ds = d.compute(i, k)	Computes the descriptors of keypoints k over I
dd = AKAZE_create()	Instantiates the AKAZE detector/descriptor
m = BFMatcher.create(NORM_HAMMING, crossCheck=True)	Instantiates a brute-force matcher, with x-checking, and Hamming distance
ms = m.match(ds_l, ds_r)	Matches the left and right descriptors
i_m = drawMatches(i_l, k_l, i_r, k_r, ms, None)	Draws matches from the left keypoints k_l on left image I_l to right I_r , using matches ms

Detection

ccs = matchTemplate(i, t, TM_CCORR_NORMED)	Matches template T to image I (normalized X-correl)
m, M, m_l, M_l = minMaxLoc(ccs)	Min, max values and respective coordinates in ccs
c = CascadeClassifier()	Creates an instance of an “empty” cascade classifier
r = c.load("file.xml")	Loads a pre-trained model from file; r is True/False
objs = c.detectMultiScale(i)	Returns 1 tuple (x, y, w, h) per detected object

Motion and Tracking

pts = goodFeaturesToTrack(i, 100, 0.5, 10)	Returns 100 Shi-Tomasi corners with, at least, 0.5 quality, and 10 pixels away from each other
pts1, st, e = calcOpticalFlowPyrLK(i0, i1, pts0, None)	New positions of pts from estimated optical flow between I_0 and I_1 ; st [<i>i</i>] is 1 if flow for point i was found, or 0 otherwise
t = TrackerCSRT_create()	Instantiates the CSRT tracker
r = t.init(f, bbox)	Initializes tracker with frame and bounding box
r, bbox = t.update(f)	Returns new bounding box, given next frame

Drawing on the image

line(i,(x0, y0),(x1, y1), (b, g, r), t)	Line
rectangle(i, (x0, y0), (x1, y1), (b, g, r), t)	Rectangle
circle(i,(x0, y0), radius, (b, g, r), t)	Circle
polylines(i,[pts], True, (b, g, r), t)	Closed (True) polygon (pts is array of points)
putText(i, "Hi", (x,y), FONT_HERSHEY_SIMPLEX, 1, (r,g,b), 2, LINE_AA)	Writes “Hi” at (x,y) , font size=1, thickness=2

Parameters

(x0, y0)	Origin/Start/Top left corner (note that it’s not (row,column))
(x1, y1)	End/Bottom right corner
(b, g, r)	Line color (uint8)
t	Line thickness (fills, if negative)

Calibration and Stereo

r, crns = findChessboardCorners(i, (n_x,n_y))	2D coords of detected corners; i is gray; r is the status; (n_x , n_y) is size of calib target
crnrs = cornerSubPix(i, crns, (5,5), (-1,-1), crit)	Improves coordinates with sub-pixel accuracy
r, K, D, ExRs, ExTs = calibrateCamera(crnrs_3D, crns_2D, i.shape[:2], None, None)	Calculates intrinsics (inc. distortion coeffs), & extrinsics (i.e. 1 R+T per target view); crns_3D contains 1 array of 3D corner coords p/target view; crns_2D contains the respective arrays of 2D corner coordinates (i.e. 1 crns p/target view)
drawChessboardCorners(i, (n_x, n_y), crns, r)	Draws corners on I (may be color); r is status from corner detection
u = undistort(i, K, D)	Undistorts I using the intrinsics
s = StereoSGBM_create(minDisparity = 0, numDisparities = 32, blockSize = 11)	Instantiates Semi-Global Block Matching method
s = StereoBM_create(32, 11)	Instantiates a simpler block matching method
d = s.compute(i_L, i_R)	Computes disparity map (α^{-1} depth map)

Termination criteria (used in e.g. K-Means, Camera calibration)

crit = (TERM_CRITERIA_MAX_ITER, 20, 0)	Stops after 20 iterations
crit = (TERM_CRITERIA_EPS, 0, 1.0)	Stop if “movement” is less than 1.0
crit = (TERM_CRITERIA_MAX_ITER TERM_CRITERIA_EPS, 20, 1.0)	Stops whatever happens first

Useful stuff

Numpy (np.)

m = mean(i)	Mean/average of array I
m = average(i, weights)	Weighted mean/average of array I
v = var(i)	Variance of array/image I
s = std(i)	Standard deviation of array/image I
h,b = histogram(i.ravel(),256,[0,256])	numpy histogram also returns the bins b
i = clip(i, 0, 255)	numpy ’s saturation/clamping function
i = i.astype(np.float32)	Converts the image type to float32 (vs. float32)
x, _, _ = linalg.lstsq(A, b)	Solves the least squares problem $\frac{1}{2}\ Ax - b\ ^2$
i = hstack((i1, i2))	Merges I_1 and I_2 side-by-side
i = vstack((i1, i2))	Merges I_1 above I_2
i = fliplr(i)	Flips image left-right
i = flipud(i)	Flips image up-down
i = pad(i, ((1, 1), (3, 3)), 'reflect')	Alternative to copyMakeBorder (also top, bottom, left, right)
idx = argmax(i)	Linear index of maximum in I (i.e. index of flattened I)
r, c = unravel_index(idx, i.shape)	2D coordinate of the index with respect to shape of i
b = any(M > 5)	Returns True if any element in array M is greater than 5
b = all(M > 5)	Returns True if all elements in array M are greater than 5
rows, cols = where(M > 5)	Returns indices of the rows and cols where elems in M are ≥ 5
coords = list(zip(rows, cols))	Creates a list with the elements of rows and cols paired
M_inv = linalg.inv(M)	Inverse of M
rad = deg2rad(deg)	Converts degrees into radians

Matplotlib.pyplot (plt.)

imshow(i, cmap="gray", vmin=0, vmax=255)	matplotlib ’s imshow preventing auto-normalization
quiver(xx, yy, i_x, -i_y, color="green")	Plots the gradient direction at positions xx , yy
savefig("name.png")	Saves the plot as an image

Copyright © 2019 António Anjos (Rev: 2020-10-09)
Most up-to-date version: <https://github.com/a-anjos/python-opencv>